

154B Discussion 6

February 15th, 2023

Goals

- Assignment 4
 - Multi-cycle memory components
 - Caching
- Caches
- Week 6 Quiz.

Logistics

- Assignment 3 due on Friday.
- Assignment 4 is going to be released sometime this week.

Assignment 4

- More complete system compared to previous assignments
 - 5-stage pipelined CPU, the same one as in assignment 3.
 - Memory request no longer takes one cycle to complete.
 - One level of memory cache, with at least 3-cycle latency.
 - RAM device now has 30-cycle latency.

Split L1 cache: L1 Instruction Cache & L1 Data Cache
(L1I) (L1D)

L₁I and L₁D are similar:

each of them has 32 entries, 4-way associative, block size = 8 bytes, LRU replacement policy.

Assignment 4: Motivation

- Multi-cycle memory component: a memory component that takes more than one CPU cycle to process a memory request.
- Expecting a memory request to be complete within 1 cycle is unrealistic.
- Real systems have caches.

Assignment 4: Motivation

Mem Hierarchy	AMD EPYC 7742 DDR4-3200 (ns @ 3.4GHz)	AMD EPYC 7601 DDR4-2400 (ns @ 3.2GHz)	Intel Xeon 8280 DDR-2666 (ns @ 2.7GHz)
L1 Cache	32KB	32KB	32KB
	4 cycles 1.18ns	4 cycles 1.25ns	4 cycles 1.48ns
L2 Cache	512KB	512KB	1024KB
	13 cycles 3.86ns	12 cycles 3.76ns	14 cycles 5.18ns
L3 Cache	16MB / CCX (4C) 256MB Total	16MB / CCX (4C) 64MB Total	38.5MB / (28C) Shared
	~34 cycles (avg) ~10.27 ns		~46 cycles (avg) ~17.5ns
DRAM	~122ns (NPS1)		
128MB Full Random	~113ns (NPS4)	~116ns	~89ns
DRAM	~134ns (NPS1)		~109ns

<https://www.anandtech.com/show/14694/amd-rome-epyc-2nd-gen/7>

Assignment 4

Part 1: Re-implement the hazard detection unit.

Part 2: Collecting data.

Part 3: Performance Analysis: Explain why there's a speedup or a slowdown.

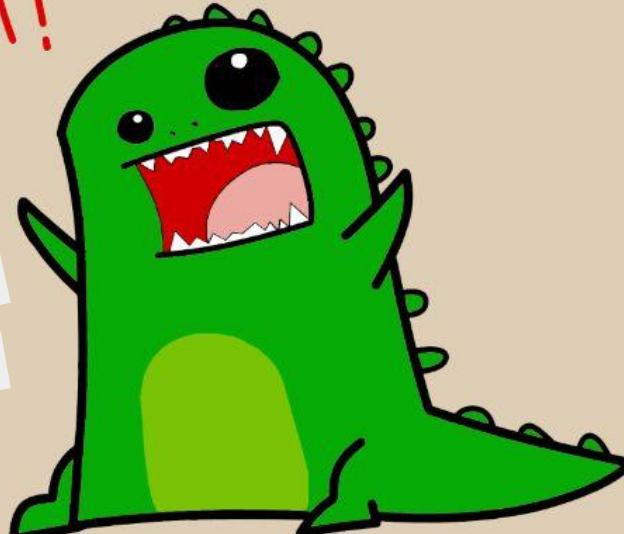
Start early! Correctness of part 2 & 3 are dependent on how reasonable part 1 is implemented.

Assignment 4: Re-implement the hazard detection unit

RAWR!

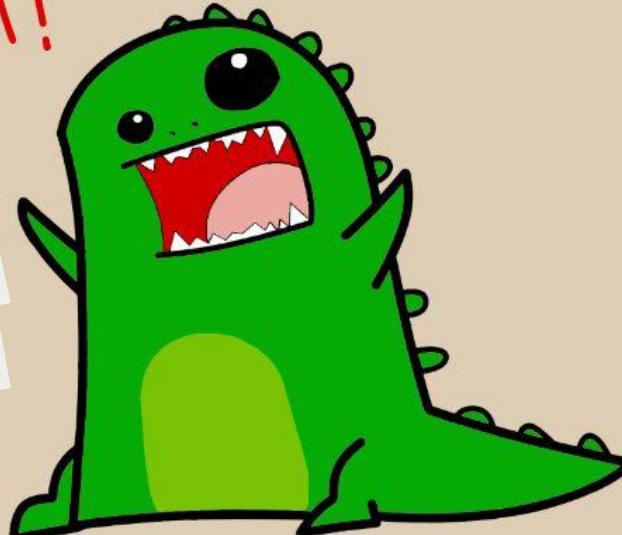
it
means

new data hazard
in DINO CPU



RAWR!

~~'t means
new data hazard
in DINO CPU~~



Assignment 4: Re-implement the hazard detection unit

- Data hazards,
 - We're not re-ordering the instructions so data hazards are still the same as the ones in assignment 3.
 - Still have RAW hazards (resolved by forwarding) and load-to-use hazards (resolved by stalling).
 - Don't need to modify forwarding unit.

Assignment 4: Re-implement the hazard detection unit

- Control hazards,
 - Still have to take care of mispredictions, i.e. jumps and taken branches.

Assignment 4: Re-implement the hazard detection unit

of

- There's no new type hazard. However, now that timing involves, we have to take care control hazards quite differently.

Assignment 4: Re-implement the hazard detection unit

- There's no new type hazard. However, now that timing involves, taking care of control hazards are more involved.

Assignment 4: The new memory interface

- Both instruction memory and data memory interfaces now have additional inputs and outputs,
 - valid (CPU inputs to memory): 1 if CPU sends the request, 0 otherwise.
 - good (Output from memory): the response is valid.
 - ready (Output from memory): the memory is ready to receive memory request.

Instruction Memory Interface

valid (Input)

address (Input)

instruction (Output)

good (Output)

ready (Output)

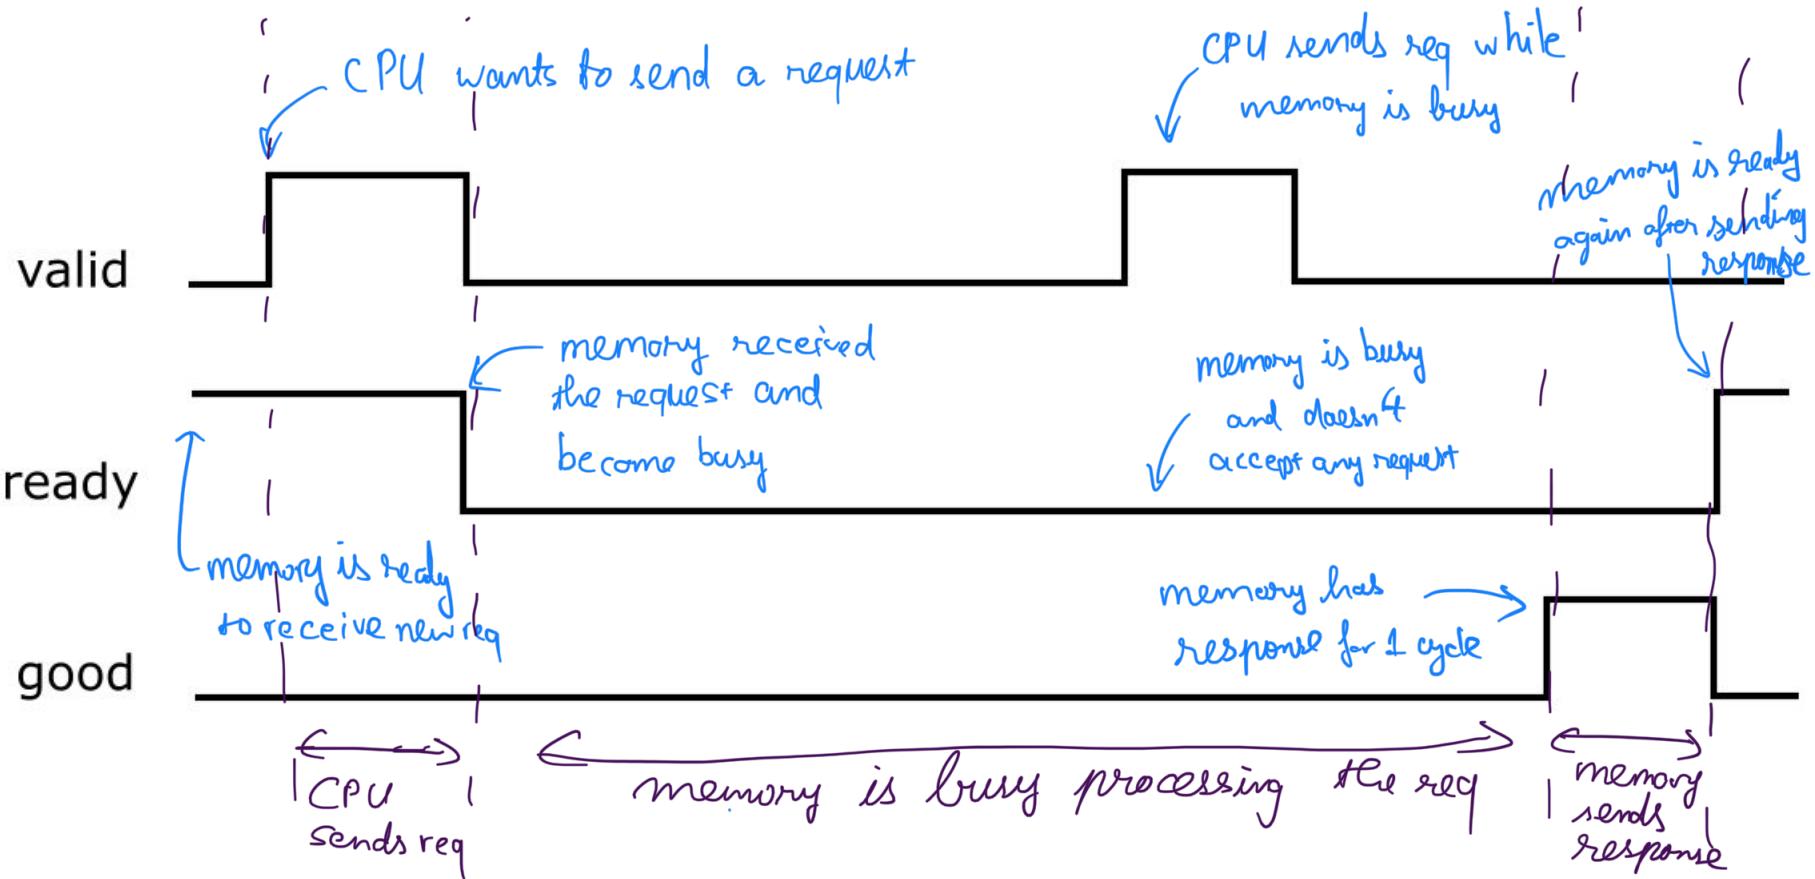
Assignment 4: The new memory interface

- The contract between the memory interface and the CPU,
 - The CPU only sends a request when valid is true and Memory is ready.
 - The memory only processes one request at a time.
 - If the memory receives a request in the current cycle, it will start processing the request in the next cycle.
 - If the memory is not ready, then even if the CPU sets valid to 1, the memory does not receive the request.
 - The memory only guarantees the response to be correct when and only when the `good` signal is 1. It means, if the `good` signal is 0, the response might contain garbage.
 - The `good` signal is only set to one for 1 cycle. It means, the response is only guaranteed to be correct in that cycle.

Assignment 4: The new memory interface

- Note: there's no guarantee in terms of memory request latency, and there's no upper bound for the latency of processing one request.
- Should rely on the good/ready signals rather than stalling the pipeline for a fixed number of cycles.

Assignment 4: The new memory interface



Assignment 4: Where are the valid instructions

- Recall: we use IF_ID, ID_EX, EX_MEM, and MEM_WB stage registers to store the instructions.
- The output of IF_ID contains the instruction in ID stage.
- The output of ID_EX contains the instruction in EX stage.
- The output of EX_MEM contains the instruction in MEM stage.
- The output of MEM_WB contains the instruction in WB stage.
- Where is the instruction in the IF stage?
 - In assignment 3, the instruction in the IF stage is in the instruction memory!
 - The instruction in the IF stage does not necessarily enter the pipeline!
- Only when the instruction reaches ID stage, it enters the pipeline.

Assignment 4: Where are the valid instructions

- Now that the output of the instruction memory is not always a correct instruction, we don't want to advance the instruction in the IF stage to the ID stage if the response from instruction memory is not "good".
- How does that look like? It's quite hard to keep the pipeline busy!

	IF	ID	EX	MEM	WB	
cycle 15, imem.good = 0	err	add	---	---	---	
cycle 16, imem.good = 0	err	???	add	---	---	
cycle 17, imem.good = 0	err	???	???	nop	add	---
cycle 18, imem.good = 1	ld	???	???	nop	add	---
cycle 18, imem.good = 0	err	ld	???	???	nop	add

flush

err : garbage

Assignment 4: Re-implement the hazard detection unit

- HazardUnit has the power of stalling/flushing the pipeline. We can use it to decide when to advance an instruction to the pipeline.
- Other issues involving timing will be discussed in later discussions.

Assignment 4: Benchmarks

- There are four benchmarks

stream-64-stride-1.riscv

stream-64-stride-4.riscv

stream-64-stride-1-noverify.riscv

stream-64-stride-4-noverify.riscv

The benchmark looks like

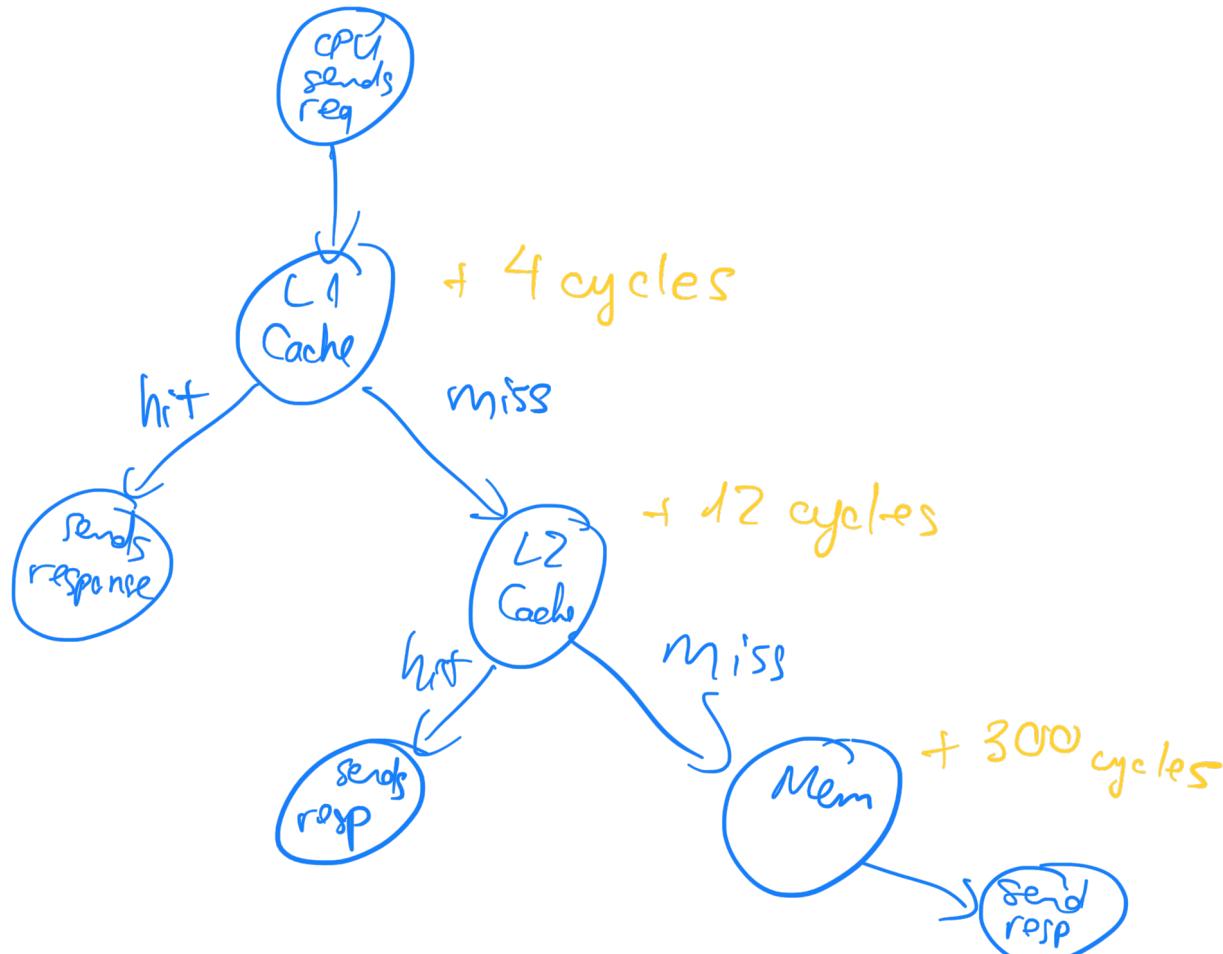
```
||| Uint16_t *src, *dst;  
||| for (size_t iter = 0; iter < 2; iter++)  
|||     for (size_t idx = 0; idx < 64; idx++)  
|||         dst[idx] = src[idx]
```

- The latter two (ones without verify functions) are used for performance evaluation.
- The first two are used for testing to make sure that the data when running the latter two are reasonable.

```
||| Uint16_t *src, *dst;  
||| for (size_t iter = 0; iter < 2; iter++)  
|||     for (size_t idx = 0; idx < 256; idx += 4)  
|||         dst[idx] = src[idx]
```

Caches

- Diagram



Caches: AMAT (average memory access time)

$$\text{AMAT} = \underset{(4 \text{ cycles})}{\text{L1 access time}} + \underset{\text{L1 miss rate}}{\text{L1 miss rate}} \times \underset{\text{L1 miss time}}{\text{L1 miss time}}$$

$$\text{where L1 miss time} = \underset{(12 \text{ cycles})}{\text{L2 access time}} + \underset{\text{L2 miss rate}}{\text{L2 miss rate}} \times \underset{\text{L2 miss time}}{\text{L2 miss time}}$$

$$\text{where L2 miss time} = \underset{(300 \text{ cycles})}{\text{main memory access time}}$$

Generalization with N-level cache:

$$\text{AMAT} = \underset{\text{L1 access time}}{\text{L1 access time}} + \underset{\text{L1 miss rate}}{\text{L1 miss rate}} \times \underset{\text{L1 miss time}}{\text{L1 miss time}}$$

$$\text{where } \left\{ \begin{array}{l} \underset{\text{L}_j \text{ miss time}}{\text{L}_j \text{ miss time}} = \underset{\text{L}_{(j+1)} \text{ access time}}{\text{L}_{(j+1)} \text{ access time}} + \underset{\text{L}_{(j+1)} \text{ miss rate}}{\text{L}_{(j+1)} \text{ miss rate}} \times \underset{\text{L}_{(j+1)} \text{ miss time}}{\text{L}_{(j+1)} \text{ miss time}} \\ \text{if } \text{L}_j \text{ is not the last level cache (LLC)} \end{array} \right.$$

$$\left. \begin{array}{l} \underset{\text{L}_j \text{ miss time}}{\text{L}_j \text{ miss time}} = \underset{\text{main memory access time}}{\text{main memory access time}} \text{ if } \text{L}_j \text{ is LLC} \end{array} \right.$$

$$\# \text{ entries} = \# \text{ rows} \times \# \text{ ways}$$

Caches: Types

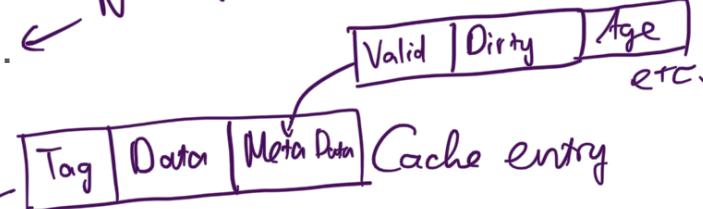
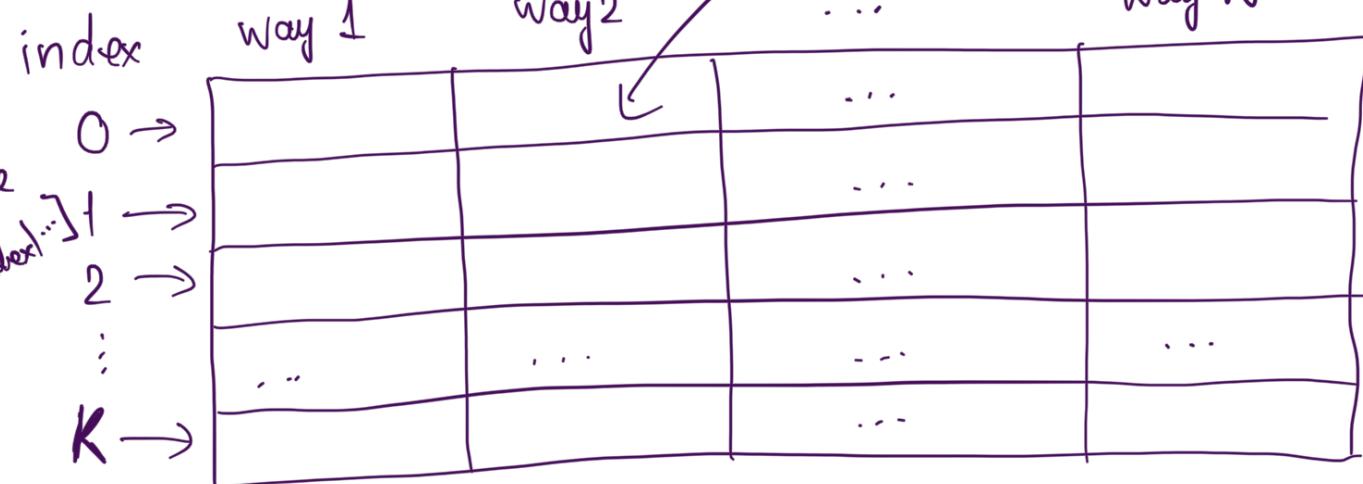
$$\rightarrow \# \text{ rows} = \frac{\# \text{ entries}}{\# \text{ ways}}$$

(for fully associative cache, $\# \text{ rows} = 1$, for others, $\# \text{ rows} = \# \text{ different indices}$)

- Direct mapped (one entry per possible index).
- N-way associative (N entries per possible index).
- Fully associative (no index). ← 1 row

Cache organization:

each cache entry must have enough information to uniquely identify the address of the data.
 \hookrightarrow address = [Tag | Index]



Caches: Tag bits, Index bits, Offset bits

Address = [Tag bits | Index bits | Offset bits]

offset bits = \log_2 (cache block size in bytes)

index bits = \log_2 (# rows)

tag bits = # address bits - # offset bits - # index bits

↳ a cache entry must have tag to recover the tag bits

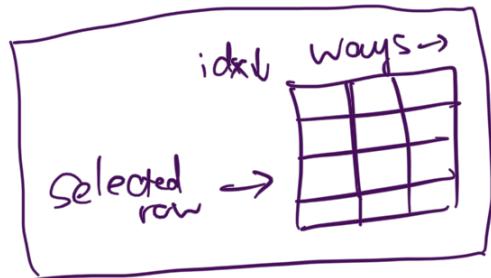
- the index bits can be recover from the index of the entry in the cache table (i.e., which row)
⇒ we know whether the cache block contains the requested address

Caches: Cache Block Size

- . An entry of a cache holds <cache-block-size> amount of data.
 - ↳ in real systems: ~ 64 bytes / block
 - ↳ in DINO CPU : ~ 8 bytes / block
- . Cache block size determines the number of offset bits = \log_2 (Cache block size in bytes)
- . E.g., req address = 0x104, block size = 64 bytes \rightarrow 6 bits for offset
 - ↳ the cache block will contain the data of address in the range: 0004 0000 0000 \rightarrow 0004 0111 1111
 $(0x400 \rightarrow 0x17F)$

~~Quiz 6~~

Where to place a new cache block?



new cache block (Data + address)

select the row according to the index bits

Compare the tags of entries in that row

↓ if there's no identical tag

is there an entry in that row
that is not valid?

if there's an
identical tag
write to that
block

yes

write to the
invalid entry

no

evict an entry from
that row according to
(replacement - policy)

↪ write to the evicted entry

~~Quiz 6~~

LRU replacement policy : the cache will evict the entry with the oldest age .
↑ least recently used

- When the age of an entry is updated?
 - upon a memory access (read from / write to that block)

Quiz 6

, 1920×1080 pixels,

a. 1 bit / pixel $\rightarrow \frac{1920 \times 1080}{8}$ bytes

b. 3 bytes / pixel $\rightarrow 1920 \times 1080 \times 3$ bytes

c. 90Hz \rightarrow transferring data 90 times / second

\rightarrow bandwidth = $1920 \times 1080 \times 3 \times 90$ bytes / second
 \rightarrow MiB / second

d. Watts = Joules / second ; $2 \times 10^{-9} \text{ J/bit} \times \text{bandwidth (in bits/second)}$
(round to the nearest integer)